

SF Bay Area
206 954 3001
linkedin.com/in/abhishekpatnia

Abhishek Patnia

SKILLS

Transformer Fundamentals: Attention mechanisms, Positional Encoding types, MoE, KV Cache; **Agentic AI & LLM Adaptation:** agents, multi-agent architectures, function calling, customer-level optimization via memory, automated LLM-based evaluations, SFT, RFT; **Platforms & Tools:** OpenAI, Huggingface, Predibase, Pytorch, Python, Pyspark on Databricks; **Classical ML:** Deep Learning, Supervised Classification \w Noisy Data; **Working Knowledge:** GPU Fundamentals, CUDA, Triton, Model quantization and distillation, vLLM, AWS, Docker, K8s, Concurrent Programming

EXPERIENCE

Nubank, SF Bay Area – Senior Staff ML Engineer/Scientist

March 2024 – PRESENT

- Technical leader on the Capabilities team, integrating skills into Nubank AI using agents. For example: [Pix Transfer Using AI](#).
- Defined the roadmap and best practices for scalable skill design and integration using agents.
- Leverage synthetic and historic datasets to iterate on agent behavior and de-risk new skill launches.
- Proved that open-weight SLMs like Phi-4-mini-instruct can power these agents. Used LoRA + GRPO fine-tuning of SLMs using existing log-traces to improve their tool-calling and instruction following performance.

Tinder, Los Angeles – Staff ML Engineer/Scientist

April 2019 – Jan 2023

- Founded and scaled the Trust & Safety ML team to 9 engineers & analysts.
- Designed multilingual text/image classification pipelines with Transformers & ConvNeXT.
- Established adversarial pre-processing pipelines and metadata embeddings to improve robustness.
- Built real-time inference pipelines (TFLite, TensorRT, Triton) deployed via Kubernetes.
- Partnered with product and risk analysts to invent KPIs and close coverage gaps across violations.

Amazon, Seattle – Senior Applied Scientist

April 2015 – April 2019

- Technical lead for the query understanding science team.
- Designed and launched Amazon's first RNN mapping query → shopping intent, powering shopping experiences at scale.
- Established deep learning training & deployment best practices, introducing GPU-based training infra with AWS Batch.
- Collaborated with infra teams to create a Python-based inference framework, ensuring seamless deployment of models across orgs.

Amazon, Seattle – Senior Software Development Engineer

June 2011 – March 2015

- Founding engineer of Kindle X-Ray, building algorithms to identify characters, passages, and images in books.
- Developed entity resolution algorithms to unify aliases (e.g., “Mr. Potter” = “Harry Potter”)
- Created large-scale test harnesses: every algorithm change was validated across thousands of books before release.
- Built the Kindle N-gram Corpus (MapReduce + DynamoDB), foundational for topic identification across the Kindle library.

ADVISING

AI Startups, SF Bay Area – Advisor

MONTH 2023 – Current

Advisor to early and growth-stage startups through the Tola Capital Advisor Program, meeting with founders and leadership teams to shape their LLM strategy, architecture, and customer-facing features. Provide customer-centric perspectives on how generative AI can drive adoption, reduce costs, and differentiate products. Partnered with companies like Martian and Instawork, helping them design and scale production-ready LLM solutions. Guidance spans agent design and orchestration, fine-tuning open-weight models with LoRA, and inference optimization & evaluation.

EDUCATION

University of Southern California – MS Computer Science

2009 – 2011, Los Angeles

Pune Institute of Computer Technology – B.E. Computer Science

2003 – 2007, India